

# TuneTrace: Semantic Music Recommendation and Intent Modeling

Bridging real-time user intent and dense catalog discovery via decoupled recurrent embeddings, self-attention, and multi-tier resilience.

**Candidate:** Agrannya Singh | Register No: [23BCE0965](#)

**School:** School of Computer Science & Engineering (SCOPE), VIT Vellore

**Implemented Stack:** PyTorch, Sentence-BERT, DuckDB, FAISS IVFPQ, FastAPI, Supabase pgvector

## Executive Definition

**TuneTrace** is an AI-driven music recommendation engine and sequential modeling sandbox designed to eliminate the two greatest failure modes of modern streaming systems: **recommendation inertia** (slow mood adaptation) and the **cold-start barrier**.

### 1. Instant Adaptation

Dynamically weights the user's active session history using exponential recency decay ( $\lambda = 1.0$ ), adapting within a single song transition.

### 2. Zero-Shot Discovery

Projects acoustic metadata into continuous 384-dimensional SBERT space, instantly placing new songs without needing listening history.

### 3. 100% Resilience

Guarantees non-empty recommendations via a 4-tier fallback hierarchy (Semantic  $\rightarrow$  Heuristic  $\rightarrow$  YouTube  $\rightarrow$  Global Charts).

### Dilemma A: Intra-Session Intent Pivot


- In a 15-minute session, a user shifts mood (e.g., from high-BPM workout to ambient study lo-fi).
- Collaborative Filtering (CF) profiles [5] rely on long-term interaction matrices that take **hours or days to decay**.
- **Failure Mode:** The user is trapped in an algorithmic echo chamber of unwanted upbeat music.

### Dilemma B: The Cold-Start Metadata Void

- New users and indie tracks have zero historical interaction entries in the interaction matrix  $\mathbf{R}$ .
- Matrix Factorization cannot compute dot products for unindexed IDs ( $\hat{r}_{ui} = \mathbf{p}_u^T \mathbf{q}_i$ ).
- **Failure Mode:** Long-tail tracks are never recommended, while new users experience completely blank slates.

### TuneTrace Solution

Replace ID-based matrix factorization with **continuous semantic embeddings [2]** and an **in-memory recency-decay profiler**.



## Behavioral Analysis

- $t = 0 \rightarrow 6$  **min**: User consumes workout music (Energy: 85).
- $t = 6.1$  **min**: Sudden switch to focus music (Energy: 18).
- **Traditional CF**: Decays linearly over hours, staying above 54 BPM at 15 min.
- **TuneTrace ( $\lambda = 1.0$ )**: Exponential decay rapidly stabilizes on new intent within **48 seconds**.

## 1. Ingest & Enrich

Extracts raw tracks from YAMDA-50M dataset. Enriches sparse metadata via YouTube API and encodes into 384-d dense vectors.

## 2. Recency Profile

Aggregates recent interactions in FastAPI memory using exponential temporal decay ( $\lambda = 1.0$ ), building a dynamic taste vector  $\mathbf{U}(t)$ .

## 3. Dual-Engine Retrieval

Performs  $< 45\text{ms}$  vector cosine search on `pgvector`, complemented by our PyTorch GRU v3 Heavy sequential model.

## 4. Resilient Filter

Applies 24h Redis deduplication and 20% diversity sampling; falls back deterministically to genre/global trends if matches equal zero.

## Engineering Guarantee

Sub-85ms total request latency, zero cold-start failures, and 100% request fulfillment rate.

## The Behavioral Limitation in YAMDA

- YAMDA dataset [6] supplies 50M rows of raw interactions:  
`<user_id,item_id,played_ratio,timestamp>`
- Completely lacks track titles, lyrics, and acoustic tags.
- Numeric IDs cannot directly resolve semantic search intents like *"late night driving beats"*.

## TuneTrace Metadata Enrichment Pipeline

- **YouTube Data API V3 Bridge:** Harvests un-enriched track IDs in automated batches of 50.
- **Ontology Filtering:** Scans snippet tags against our `KNOWN_MUSIC_GENRES` taxonomy.
- **Database Binding:** Saves structured tags (`enriched="V2"`), enabling 384-d SBERT vectorization.

aggregation

## Mathematical Formulation

$$\mathbf{v}_i = \text{SentenceTransformer}(\text{Metadata}_i) \in \mathbb{R}^{384} \quad [2]$$

$$w_i = \frac{\exp(-\lambda \cdot \frac{n-i}{n})}{\sum_{j=1}^n \exp(-\lambda \cdot \frac{n-j}{n})}, \quad \mathbf{U}(t) = \frac{\sum_{i=1}^n w_i \mathbf{v}_i}{\left\| \sum_{i=1}^n w_i \mathbf{v}_i \right\|_2}$$


$$D_C(\mathbf{U}(t), \mathbf{w}) = 1 - \frac{\mathbf{U}(t) \cdot \mathbf{w}}{\|\mathbf{U}(t)\|_2 \|\mathbf{w}\|_2}$$

## Parameter Definitions


- $\mathbf{v}_i \in \mathbb{R}^{384}$ : L2-normalized embedding of track  $i$  via all-MiniLM-L6-v2.
- $n \in \mathbb{N}$ : Total number of items in active user session history.
- $i \in \{1, \dots, n\}$ : Chronological event index ( $i = n$  is the latest track).
- $\lambda = 1.0$ : Exponential recency decay sensitivity parameter.
- $w_i \in (0, 1)$ : Normalized attention weight assigned to track  $i$  ( $\sum w_i = 1$ ).
- $\mathbf{U}(t) \in \mathbb{R}^{384}$ : Dynamic unit-norm user taste profile vector.
- $\mathbf{w} \in \mathbb{R}^{384}$ : Candidate track embedding in catalog.
- $D_C \in [0, 2]$ : Cosine distance metric (threshold  $\leq 0.28$ ).

# System Architecture Schematic

Three-tier decoupled microservice







## Stage 1: Ingestion (**DuckDB** [7])

Parquets streamed with 8 GB RAM cap; SQLite stores relations while NumPy Memmap hosts the L2-normalized embedding matrix.

## Stage 2: FAISS Indexing (IVFPQ [4])

Product Quantization compresses  $\geq 50\text{k}$  vectors for 90% RAM reduction with sub-millisecond retrieval.

### Stage 3: Deep GRU Training (**GRU [1]** / **PyTorch [8]**)

4-layer GRU with Self-Attention pooling trained via InfoNCE Contrastive + Intent Cross-Entropy loss.

## Stage 4: Recommendation Generation

Calibrates FAISS top-5 ANN queries:  $\text{score} \times (0.5 + 0.5 \cdot P(\text{like}))$ .

### ETL Architecture

- **Disk-Backed Execution:** Enforced an **8 GB memory ceiling** to prevent OOM errors.
- **Vectorized Joins:** Joins `embeddings.parquet` and `listens.parquet` directly on-disk.
- **Noise Filtering:** Excluded passive listening records where `played_ratio_pct < 0.15`.

### Measured Benchmarks

- **Catalog Size:** **156,275 tracks** extracted and indexed.
- **User Sequences:** **427 users** parsed with complete chronological logs.
- **Ingestion Time:** Completed in **118 seconds**.
- **RAM Footprint:** Peak execution memory remained below **450 MB**.

### Mathematical Formulation

$$\mathbf{x} = [\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}] \in \mathbb{R}^D$$

$$q(\mathbf{x}) = [\mathbf{c}_{i_1}^{(1)}, \mathbf{c}_{i_2}^{(2)}, \dots, \mathbf{c}_{i_m}^{(m)}]$$

$$d_{ADC}^2(\mathbf{U}, q(\mathbf{w})) = \sum_{j=1}^m \|\mathbf{U}^{(j)} - \mathbf{c}_{i_j}^{(j)}\|_2^2$$

**RAM Saving:** 32 B vs 1024 B ( $\approx 90\%$  reduction).

### Parameter Definitions

- $\mathbf{x} \in \mathbb{R}^D$ : High-dimensional track embedding ( $D = 256$ ).
- $m = 32$ : Number of orthogonal sub-vector slices ( $d = D/m = 8$ ).
- $\mathbf{x}^{(j)} \in \mathbb{R}^8$ :  $j$ -th sub-vector component of track  $\mathbf{x}$ .
- $\mathcal{C}^{(j)}$ : Centroid codebook ( $k^* = 256$ , 8-bit quantization per slice).
- $\mathbf{c}_{i_j}^{(j)}$ : Nearest centroid vector in codebook  $j$  ( $i_j \in [0, 255]$ ).
- $q(\mathbf{x})$ : Quantized index represented as a 32-byte code.
- $\mathbf{U} \in \mathbb{R}^{256}$ : Unquantized user query profile vector.
- $d_{ADC}^2$ : Asymmetric Distance Computation via Look-Up Tables.

## Recurrent Gating Dynamics

$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{U}_z \mathbf{h}_{t-1} + \mathbf{b}_z)$$

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{U}_r \mathbf{h}_{t-1} + \mathbf{b}_r)$$

$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W}_h \mathbf{x}_t + \mathbf{U}_h (\mathbf{r}_t \odot \mathbf{h}_{t-1}) + \mathbf{b}_h)$$

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t$$

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left( \frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}} \right) \mathbf{V}$$

## Parameter Definitions

- $\mathbf{x}_t \in \mathbb{R}^{256}$ : Input track embedding at timestep  $t \in [1, 20]$ .
- $\mathbf{h}_t \in \mathbb{R}^{512}$ : Hidden state vector across 4 GRU layers.
- $\mathbf{z}_t \in (0, 1)^{512}$ : Update gate vector (controls memory retention).
- $\mathbf{r}_t \in (0, 1)^{512}$ : Reset gate vector (controls history forgetfulness).
- $\tilde{\mathbf{h}}_t$ : Candidate hidden state activation via tanh.
- $\mathbf{W}, \mathbf{U}, \mathbf{b}$ : Learnable input, recurrent, and bias matrices.
- $\odot$ : Element-wise (Hadamard) tensor multiplication.
- $\mathbf{Q}, \mathbf{K}, \mathbf{V}$ : Query ( $\mathbf{h}_{20}$ ), Key ( $\mathbf{H}$ ), Value ( $\mathbf{H}$ ) projections ( $d_k = 128$ ).

## Loss Objectives

$$\mathcal{L}_{\text{InfoNCE}} = -\log \frac{\exp\left(\frac{\mathbf{p} \cdot \mathbf{y}^+}{\tau}\right)}{\exp\left(\frac{\mathbf{p} \cdot \mathbf{y}^+}{\tau}\right) + \sum_{j=1}^{B-1} \exp\left(\frac{\mathbf{p} \cdot \mathbf{y}_j^-}{\tau}\right)}$$

$$\mathcal{L}_{\text{CrossEntropy}} = -\sum_{c=1}^4 y_c \log \hat{y}_c$$

$$\mathcal{L}_{\text{joint}} = \mathcal{L}_{\text{InfoNCE}} + 0.3 \cdot \mathcal{L}_{\text{CrossEntropy}}$$

## Parameter Definitions

- $\mathbf{p} \in \mathbb{R}^{256}$ : L2-normalized predicted embedding from projection head.
- $\mathbf{y}^+ \in \mathbb{R}^{256}$ : Ground-truth positive next-track embedding.
- $\mathbf{y}_j^- \in \mathbb{R}^{256}$ : In-batch negative samples ( $B - 1$  sample vectors).
- $\tau$ : Learnable temperature parameter (initialized at  $\tau_0 = 0.07$ ).
- $B$ : Batch size ( $B = 2048$  on CUDA AMP,  $B = 512$  on CPU).
- $y_c, \hat{y}_c$ : Ground truth and predicted logits for 4 intent classes.
- 0.3: Empirical hyperparameter balancing contrastive and classification loss.

### Empirical Diagnosis of 30-Epoch Baseline

- Initial runs with 2-layer 256-dim GRU yielded:

**Epoch 30/30 | Avg Loss: 5.9199**

InfoNCE: 5.6990, Intent: 0.7362


- Root Cause:** The 256-dim backbone lacked expressive capacity for 50M interactions.
- Gradient Starvation:** Final-step pooling lost key transitions earlier in the sequence.

### Architectural Interventions Applied

- 4× Backbone Width:** Scaled hidden dimension from 256 to 512 (> 3.5M params).
- Attention Pooling:** Weights critical actions anywhere in history (e.g. LIKE vs SKIP).
- Residual Projections:** Skip connections ensure direct gradient flow to early cells.
- Cosine Annealing:** Bypasses local plateaus.

# Empirical Training Dynamics: GRU Convergence Curve

Trajectory & hyperparameter validation



## Configured Hyperparameters

| Component          | Production Value            |
|--------------------|-----------------------------|
| Backbone           | 4-Layer GRU (512-d)         |
| Attention Heads    | 4 Heads (embed = 512)       |
| Sequence Length    | $T = 20$ Interactions       |
| Batch Size         | 2048 (CUDA AMP) / 512 (CPU) |
| Base LR            | $3 \times 10^{-3}$ (AdamW)  |
| Schedulers         | 5-epoch Warmup + Cosine     |
| Temperature $\tau$ | Learnable (init 0.07)       |
| Checkpoints        | Every 15 Epochs             |

### Long-Horizon Training Requirements

- Training a 3.5M parameter network on 50M records requires extended multi-stage execution.
- Training sessions must be resilient to runtime interruptions and system preemption.
- Uncheckpointed training risks loss of progress and model state degradation.

### Engineered Auto-Resume Logic

- **15-Epoch Checkpointing:** Automatically saves full state dicts at epochs 15, 30, 45, and 60.
- **State Persistence:** Preserves:  
$$\langle \mathbf{W}_{\text{model}}, \mathbf{S}_{\text{optimizer}}, \mathbf{S}_{\text{scheduler}}, \mathbf{S}_{\text{scaler}}, \log \tau \rangle$$
- **Zero-Intervention Resume:** Auto-loads from `yamda_checkpoint.pt` and seamlessly resumes training.



| User ID  | Input History Sequence ( $T = 5$ )            | GRU Recommender Top 5 (Deep)                  | Heuristic Engine Top 5 (Recency)              |
|----------|-----------------------------------------------|-----------------------------------------------|-----------------------------------------------|
| User 100 | [3175251, 1299533, 7415847, 6870586, 4734787] | [3443475, 6060302, 4531263, 4805129, 2413614] | [4734787, 3426714, 5150945, 451318, 2077046]  |
| User 200 | [2859641, 2943786, 4117955, 5134208, 3778807] | [5811656, 4219887, 3361052, 1252472, 3778807] | [3778807, 3741007, 1346490, 5541739, 5811656] |

Heuristic Recency Engine Behavior

Directly surfaces recent **historical tracks** (e.g., track **3778807** for User 200), prioritizing immediate session mood continuity.

Deep GRU Model Behavior

Infers broader **latent relationships** (e.g., discovering track **5811656**), uncovering non-linear semantic transitions.

| Tier   | Method             | Trigger Condition & Latency Budget                                                                       |
|--------|--------------------|----------------------------------------------------------------------------------------------------------|
| Tier 1 | Semantic Match     | Preferred method; $D_C(\mathbf{U}, \mathbf{w}) \leq 0.28$ via Supabase pgvector HNSW index ( $< 45$ ms). |
| Tier 2 | Metadata Heuristic | Triggered if vector search yields 0 results; exact genre match from catalog SQLite ( $< 15$ ms).         |
| Tier 3 | Genre Trending     | Triggered if local metadata fails; live query to YouTube Data API V3 in genre ( $< 120$ ms).             |
| Tier 4 | Global Trending    | Ultimate fallback; retrieves live top global charts to eliminate empty query errors ( $< 90$ ms).        |

### Redis Deduplication Cache

24-hour sliding-window filter excludes recently suggested IDs, completely eliminating song fatigue.

### Diversity Sampler

Injects 20% exploration tracks from the top-50 pool using Gibbs sampling to shatter filter bubbles.

### Verified Request: POST /suggestions

```
{
  "user_id": "user@example.com",
  "songs": ["Shape of You - Ed Sheeran"],
  "genre": "Pop"
}
```

### Execution Lifecycle

- 1 JWT Authentication (`get_current_user`).
- 2 In-memory profile calculation (< 1.2 ms).
- 3 pgvector Cosine distance lookup (< 45 ms).
- 4 Redis 24hr deduplication filter.
- 5 Background task logging to PostgreSQL.

### Verified Response Payload

```
{
  "status": "success",
  "source": "semantic_search",
  "results": [
    {
      "title": "Perfect",
      "artist": "Ed Sheeran",
      "score": 0.89
    },
    {
      "title": "Someone You Loved",
      "artist": "Lewis Capaldi",
      "score": 0.85
    }
  ]
}
```

## Review 1 Deliverables Completed

- DuckDB ETL processing 156k tracks in 118 seconds.
- FAISS IVFPQ index achieving 90% RAM compression.
- PyTorch GRU v3 Heavy (3.5M params) architecture.
- 15-epoch auto-resume checkpointing engine.
- Live FastAPI service with pgvector and Redis cache.

## Phase 2 Roadmap

- Execute 60-epoch training to convergence.
- Deploy GRU inside FastAPI for live A/B testing.
- Index full 50M YAMDA catalog via distributed FAISS.

## Academic References (IEEE Format)

- [1] B. Hidasi *et al.*, "Session-based recommendations with recurrent neural networks," in *Proc. 4th Int. Conf. Learn. Represent. (ICLR)*, 2016. [[arXiv:1511.06939](#)]
- [2] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using siamese BERT-networks," in *Proc. Conf. EMNLP*, 2019, pp. 3982–3992. [[arXiv:1908.10084](#)]
- [3] A. van den Oord, Y. Li, and O. Vinyals, "Representation learning with contrastive predictive coding," *arXiv preprint arXiv:1807.03748*, 2018. [[arXiv:1807.03748](#)]
- [4] J. Johnson, M. Douze, and H. Jégou, "Billion-scale similarity search with GPUs," *IEEE Trans. Big Data*, vol. 7, no. 3, pp. 535–547, 2021. [[IEEE/arXiv:1702.08734](#)]
- [5] M. Spišák *et al.*, "Scalable approximate nonsymmetric autoencoder for collaborative filtering," in *Proc. 17th ACM RecSys*, 2023, pp. 838–844. [[RecSys/arXiv:2308.06733](#)]
- [6] P. Loshkin, "YAMDA: Yet another music dataset," *GitHub repository / Hugging Face*, 2023. [[Online Dataset](#)]
- [7] M. Raasveldt and H. Mühleisen, "DuckDB: An embeddable analytical database," in *Proc. ACM SIGMOD*, 2019, pp. 1981–1984. [[ACM SIGMOD](#)]
- [8] A. Paszke *et al.*, "PyTorch: An imperative style, high-performance deep learning library," in *Adv. Neural Inf. Process. Syst. 32 (NeurIPS)*, 2019, pp. 8024–8035. [[NeurIPS/arXiv](#)]

**Thank You | Questions & Technical Discussion**